# **Importing Data and Important Libraries**

```
In [78]:
         import numpy as np
         import pandas as pd
         import matplotlib.pyplot as plt
         import seaborn as sns
         from statsmodels.tsa.stattools import adfuller
         import statsmodels.api as sm
         import warnings
         warnings.filterwarnings(action="ignore")
         from sklearn.metrics import (
             mean squared error as mse,
             mean absolute error as mae,
             mean absolute percentage error as mape
         from statsmodels.tsa.statespace.sarimax import SARIMAX
         import itertools
         from prophet import Prophet
         data = pd.read csv('train 1.csv')
In [2]:
         data.head()
 In [3]:
```

Out[3]:		Page	2015- 07-01	2015- 07-02	2015- 07-03	2015- 07-04	2015- 07-05	2015- 07-06	2015- 07-07	2015- 07-08	2015- 07-09	•••	2016- 12-22	2016- 12-23	2016- 12-24	2016- 12-25	2016- 12-26	
	0	2NE1_zh.wikipedia.org_all- access_spider	18.0	11.0	5.0	13.0	14.0	9.0	9.0	22.0	26.0		32.0	63.0	15.0	26.0	14.0	20.0
	1	2PM_zh.wikipedia.org_all- access_spider	11.0	14.0	15.0	18.0	11.0	13.0	22.0	11.0	10.0		17.0	42.0	28.0	15.0	9.0	30.0
	2	3C_zh.wikipedia.org_all-access_spider	1.0	0.0	1.0	1.0	0.0	4.0	0.0	3.0	4.0		3.0	1.0	1.0	7.0	4.0	4.0
	3	4minute_zh.wikipedia.org_all- access_spider	35.0	13.0	10.0	94.0	4.0	26.0	14.0	9.0	11.0		32.0	10.0	26.0	27.0	16.0	11.0
	4	52_Hz_I_Love_You_zh.wikipedia.org_all-access s	NaN		48.0	9.0	25.0	13.0	3.0	11.0								

5 rows × 551 columns

# **Feature Engineering**

```
In [4]: def parse_page_name(page_name):
            parts = page_name.split("_")
            if len(parts) >= 4:
                title = "_".join(parts[:-3])
                language = parts[-3].split('.')[0]
                access_type = parts[-2]
                access_origin = parts[-1]
                return title, language, access_type, access_origin
             return None, None, None, None
        # Apply the function to the `page_name` column and expand results into new columns
        data[["Title", "Language", "Access Type", "Access Origin"]] = data["Page"].apply(parse_page_name).apply(pd.Series)
```

```
data.head()
In [5]:
```

Out[5]:		Page	2015- 07-01	2015- 07-02	2015- 07-03	2015- 07-04			2015- 07-07	2015- 07-08		•••	2016- 12-26	2016- 12-27	2016- 12-28	2016- 12-29		
	0	2NE1_zh.wikipedia.org_all- access_spider	18.0	11.0	5.0	13.0	14.0	9.0	9.0	22.0	26.0		14.0	20.0	22.0	19.0	18.0	20.0
	1	2PM_zh.wikipedia.org_all- access_spider	11.0	14.0	15.0	18.0	11.0	13.0	22.0	11.0	10.0		9.0	30.0	52.0	45.0	26.0	20.0
	2	3C_zh.wikipedia.org_all-access_spider	1.0	0.0	1.0	1.0	0.0	4.0	0.0	3.0	4.0		4.0	4.0	6.0	3.0	4.0	17.0
	3	4minute_zh.wikipedia.org_all- access_spider	35.0	13.0	10.0	94.0	4.0	26.0	14.0	9.0	11.0		16.0	11.0	17.0	19.0	10.0	11.0
	4	52_Hz_I_Love_You_zh.wikipedia.org_all-access_s	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN		3.0	11.0	27.0	13.0	36.0	10.0

5 rows × 555 columns

[6]: d	data = data.melt(id_vars=["Page", "Title", "Language", "Access Type", "Access Origin"], var_name="Date", value_name="Value")												
[7]: d	lat	ca.head()											
[7]:		Page	Title	Language	Access Type	Access Origin	Date	Value					
0	)	2NE1_zh.wikipedia.org_all-access_spider	2NE1	zh	all-access	spider	2015-07-01	18.0					
1		2PM_zh.wikipedia.org_all-access_spider	2PM	zh	all-access	spider	2015-07-01	11.0					
2		3C_zh.wikipedia.org_all-access_spider	3C	zh	all-access	spider	2015-07-01	1.0					
3		4minute_zh.wikipedia.org_all-access_spider	4minute	zh	all-access	spider	2015-07-01	35.0					
4		52_Hz_I_Love_You_zh.wikipedia.org_all-access_s	52_Hz_I_Love_You	zh	all-access	spider	2015-07-01	NaN					

# **Checking Basic Metrics**

```
(79784650, 7)
Out[8]:
         data.size
In [9]:
         558492550
Out[9]:
         data.isna().sum()
In [10]:
                                0
         Page
Out[10]:
         Title
                                0
         Language
         Access Type
         Access Origin
                                0
         Date
         Value
                          6192931
         dtype: int64
         data.info()
In [11]:
         <class 'pandas.core.frame.DataFrame'>
         RangeIndex: 79784650 entries, 0 to 79784649
         Data columns (total 7 columns):
              Column
                             Dtype
                             ----
                             object
              Page
              Title
                            object
          1
              Language
                            object
              Access Type
                            object
              Access Origin object
          5
              Date
                             object
              Value
                            float64
         dtypes: float64(1), object(6)
         memory usage: 4.2+ GB
         data.describe()
In [12]:
```

```
      count
      7.359172e+07

      mean
      1.419986e+03

      std
      8.669325e+04

      min
      0.000000e+00

      25%
      1.900000e+01

      50%
      1.470000e+02

      75%
      6.670000e+02

      max
      6.726426e+07
```

# Some of the columns are don't have language so replacing them with unknown

```
data["Language"] = data["Language"].replace(["www", "commons"], "unknown")
          (data['Date'].min() , data['Date'].max())
In [14]:
          ('2015-07-01', '2016-12-31')
Out[14]:
In [15]:
          data['Date'].value counts()
          Date
Out[15]:
          2016-12-31
                        145063
          2015-07-01
                        145063
          2015-07-02
                        145063
          2015-07-03
                        145063
          2015-07-04
                        145063
                         . . .
          2015-07-11
                        145063
          2015-07-10
                        145063
          2015-07-09
                        145063
          2015-07-08
                        145063
          2015-07-07
                        145063
          Name: count, Length: 550, dtype: int64
```

# **Quick Analysis**

- Here the every date has 145063 columns so it means every page has data for every date in the dataset
- Filling the missing Nan using linear interpolation because filling the value with mean might create an discrepancy in the data

### Imputing Nan values using Linear Interpolation

### Renaming the Language column names

```
data['Language'].value_counts()
In [18]:
          Language
Out[18]:
                     13259400
          jа
                     11237050
          de
                     10200850
          unknown
                      9820250
          fr
                      9791100
                      9475950
          zh
                      8262100
          ru
          es
                      7737950
          Name: count, dtype: int64
          language_map = {
In [19]:
              "en": "English",
              "ja": "Japanese",
              "de": "German",
```

```
"fr": "French",
    "zh": "Chinese",
    "ru": "Russian",
    "es": "Spanish"
}
data['Language'] = data['Language_map)
```

In [20]: data.head()

Out[20]:		Page	Title	Language	Access Type	Access Origin	Date	Value
	0	2NE1_zh.wikipedia.org_all-access_spider	2NE1	Chinese	all-access	spider	2015-07-01	18.0
	1	2PM_zh.wikipedia.org_all-access_spider	2PM	Chinese	all-access	spider	2015-07-01	11.0
	2	3C_zh.wikipedia.org_all-access_spider	3C	Chinese	all-access	spider	2015-07-01	1.0
	3	4minute_zh.wikipedia.org_all-access_spider	4minute	Chinese	all-access	spider	2015-07-01	35.0
	4	52_Hz_I_Love_You_zh.wikipedia.org_all-access_s	52 Hz I Love You	Chinese	all-access	spider	2015-07-01	23.5

# **Non Graphical Analysis**

```
In [21]: col=['Access Type', 'Access Origin', 'Language']
for i in col:
    print("Value count of :",i)
    print(data[i].value_counts())
    print()
    print("-"*100)
```

```
Value count of : Access Type
Access Type
all-access
             40873250
mobile-web
             19766450
desktop
             19144950
Name: count, dtype: int64
Value count of : Access Origin
Access Origin
all-agents
             60582500
spider
             19202150
Name: count, dtype: int64
Value count of : Language
Language
English
           13259400
Japanese
           11237050
German
           10200850
French
            9791100
Chinese
         9475950
Russian
            8262100
Spanish
            7737950
Name: count, dtype: int64
```

## **Graphical Analysis**

```
In [22]: plt.figure(figsize=(7, 5))
    sns.set_theme(style="whitegrid")
    ax = sns.countplot(data=data, x="Language", palette="viridis")

# Add Labels and a title
    ax.set_title("Language Distribution", fontsize=14, fontweight='bold', pad=20)
    ax.set_xlabel("Language", fontsize=12)
    ax.set_ylabel("Count", fontsize=12)


# Improve readability of x-axis
    plt.xticks(rotation=45, fontsize=10, ha="right")

for container in ax.containers:
```

```
ax.bar_label(container, fmt='%d', label_type='edge', fontsize=10, padding=2)

# Show the plot
plt.tight_layout()
plt.show()
```

# Language Distribution



# **Insights:**

- Around 17% of the pages are in English, making it the most common language.
- Japanese comes next, with about 14% of the pages.
- The remaining languages each make up roughly 12% of the total.

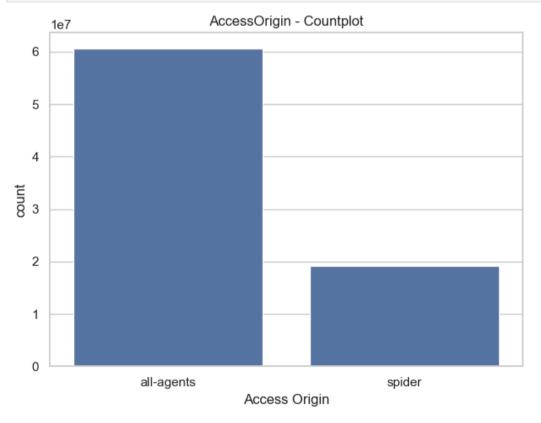
```
In [23]: fig, axes = plt.subplots(1, 2, figsize=(12, 5))

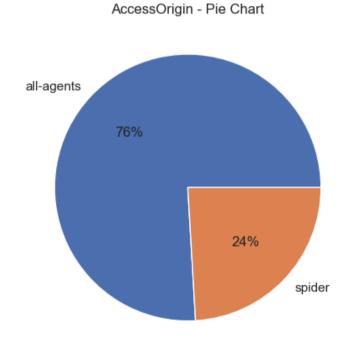
# Countplot
sns.countplot(data=data, x="Access Origin", ax=axes[0], order=data["Access Origin"].value_counts().index)
axes[0].set_title("AccessOrigin - Countplot")

# Pie chart
labels = data["Access Origin"].value_counts().index
counts = data["Access Origin"].value_counts().values


axes[1].pie(counts, labels=labels, autopct='%1.0f%%')
axes[1].set_title("AccessOrigin - Pie Chart")

plt.tight_layout()
plt.show()
```






```
In [24]: sns.countplot(y=data["Access Type"])
   plt.show()
```




# **Bivariate Analysis**

```
In [25]: plt.figure(figsize=(10,5))
    sns.countplot(data=data,x="Language",hue="Access Type")
    plt.show()
```



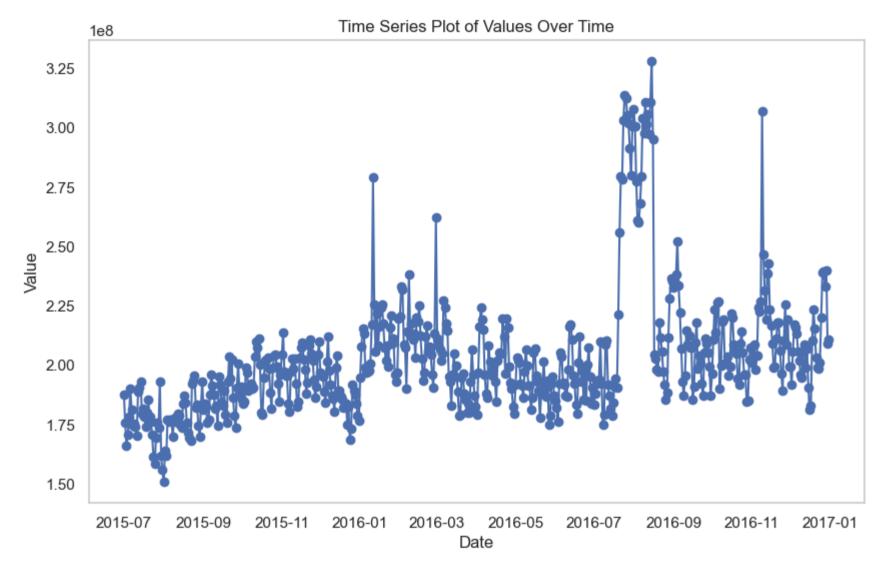
```
In [26]: plt.figure(figsize=(10,5))
    sns.countplot(data=data,x="Language",hue="Access Origin")
    plt.show()
```



In [27]: data.head()

0	u'	t		2	7		
			_			_	

	Page	Title	Language	Access Type	Access Origin	Date	Value
0	2NE1_zh.wikipedia.org_all-access_spider	2NE1	Chinese	all-access	spider	2015-07-01	18.0
1	2PM_zh.wikipedia.org_all-access_spider	2PM	Chinese	all-access	spider	2015-07-01	11.0
2	3C_zh.wikipedia.org_all-access_spider	3C	Chinese	all-access	spider	2015-07-01	1.0
3	4minute_zh.wikipedia.org_all-access_spider	4minute	Chinese	all-access	spider	2015-07-01	35.0
4	52_Hz_I_Love_You_zh.wikipedia.org_all-access_s	52_Hz_I_Love_You	Chinese	all-access	spider	2015-07-01	23.5




```
In [29]: data.set_index('Date', inplace=True)
In [30]: reshaped_df = data.pivot_table(index=data.index, columns='Language', values='Value', aggfunc='mean')
In [31]: reshaped_df.columns.name = None
In [32]: reshaped_df
```

Out[32]:		Chinese	English	French	German	Japanese	Russian	Spanish
	Date							
	2015-07-01	302.355734	4427.708334	1129.034804	790.590708	601.360077	651.452513	1209.999028
	2015-07-02	268.117246	4392.444278	502.047865	775.655888	687.995484	655.053777	1160.690263
	2015-07-03	267.042942	4166.937545	484.547462	745.781675	616.659299	609.369040	1056.775572
	2015-07-04	259.622370	4322.986667	526.252510	678.666147	776.318452	572.553909	992.093234
	2015-07-05	277.441749	4482.500642	507.830786	791.054515	750.043471	628.859743	1077.254428
	2016-12-27	379.421063	6063.200286	859.794574	1116.103251	809.202535	1003.665424	1158.402587
	2016-12-28	382.057316	5889.543575	776.484440	1063.900065	810.996280	934.810545	1205.360118
	2016-12-29	354.920570	6268.441202	754.271205	1022.153879	887.976041	900.150945	1134.710463
	2016-12-30	359.611527	5217.368425	702.872149	975.896722	959.556825	805.680469	832.056507
	2016-12-31	369.184021	5139.305521	647.605466	914.112471	1215.738192	883.976202	796.098443

550 rows × 7 columns

```
In [33]: plt.figure(figsize=(8, 6))
    sns.set_theme()
    reshaped_df.plot(ax=plt.gca())
    plt.xlabel("Date")
    plt.ylabel("Daily Average Page Views")
    plt.show()
```



# Insight

- English pages are the most visited pages follwed by Spanish
- English pages have an upward trend in terms of visits
- There is an unusual peak from mid of July to end of August 2016

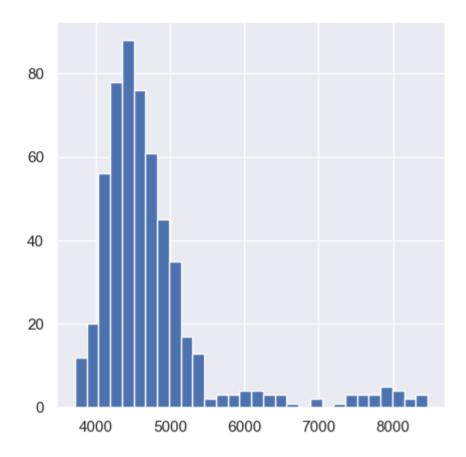
### Checking if the data is Stationary

- Null Hypothesis: Data is Not stationary
- Alternate Hypothesis: Data is stationary

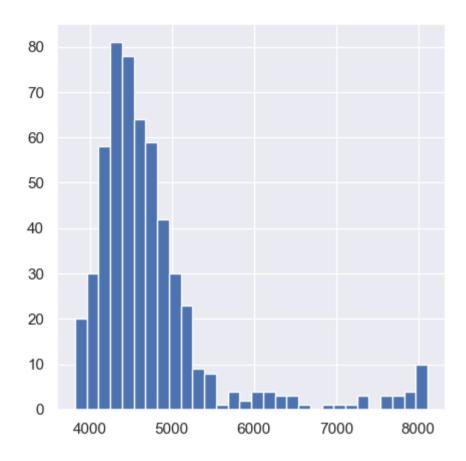
```
In [34]: def dicky fuller(data, alpha):
              result = adfuller(data)
             if result[1] < alpha:</pre>
                  print(f"The data is Stationary {result[1].round(2)}")
              else:
                  print(f"The data is non stationary {result[1].round(2)}")
In [35]: for i in reshaped_df.columns:
              print(i)
             dicky fuller(reshaped df[i], 0.05)
              print()
         Chinese
         The data is non stationary 0.22
          English
         The data is non stationary 0.06
          French
         The data is Stationary 0.03
         German
         The data is non stationary 0.1
          Japanese
         The data is non stationary 0.08
         Russian
         The data is Stationary 0.0
          Spanish
         The data is non stationary 0.05
```

Making the data Stationary for language English only

```
English = reshaped df['English']
In [36]:
         English
In [37]:
         Date
Out[37]:
         2015-07-01
                       4427.708334
         2015-07-02
                       4392.444278
         2015-07-03
                       4166.937545
         2015-07-04
                       4322.986667
         2015-07-05
                       4482.500642
                          . . .
                       6063.200286
         2016-12-27
         2016-12-28
                       5889.543575
         2016-12-29
                       6268.441202
         2016-12-30
                       5217.368425
         2016-12-31
                       5139.305521
         Name: English, Length: 550, dtype: float64
In [38]: plt.figure(figsize=(5,5))
         English.hist(bins=30);
```




In [39]: plt.figure(figsize=(5,5))
 English\_clipped = English.clip(lower=English.quantile(0.01),upper=English.quantile(0.99))
 English\_clipped.hist(bins=30);



# Decomposing the data for english Language using Triple Exponential Smoothing

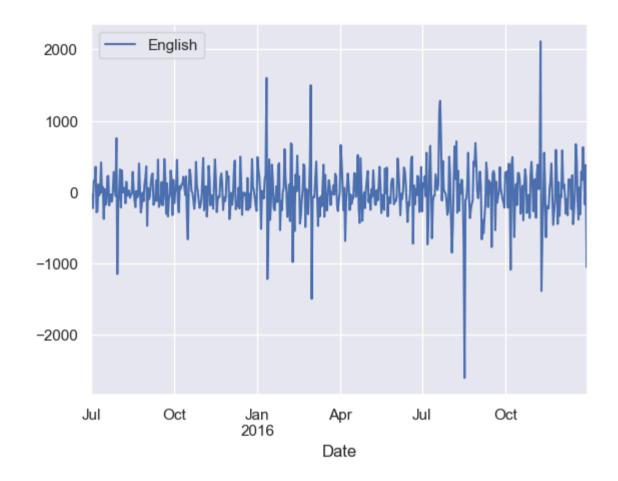
In [40]: decomp\_english = sm.tsa.seasonal\_decompose(English\_clipped)
 decomp\_english.plot();



2015-07 2015-09 2015-11 2016-01 2016-03 2016-05 2016-07 2016-09 2016-11

In [41]: English\_clipped.diff().plot()
 plt.legend()

Out[41]: <matplotlib.legend.Legend at 0x2068b6e4910>




# Checking the Stationary of the Data


```
In [42]: dicky_fuller(English_clipped.diff().dropna(), 0.05)
The data is Stationary 0.0
In [43]: Stationary_data = English_clipped.diff().dropna()
In [44]: Stationary_data
```

```
Date
Out[44]:
          2015-07-02
                         -35.264056
         2015-07-03
                        -225.506733
         2015-07-04
                        156.049122
         2015-07-05
                        159.513974
          2015-07-06
                         354.473387
          2016-12-27
                         279.345881
         2016-12-28
                        -173.656711
         2016-12-29
                         378.897627
         2016-12-30
                       -1051.072777
          2016-12-31
                         -78.062904
         Name: English, Length: 549, dtype: float64
```

# Plotting ACF and PACF plot

```
In [45]: sm.tsa.graphics.plot_acf(Stationary_data);
sm.tsa.graphics.plot_pacf(Stationary_data);
```





# ACF (Autocorrelation Function):

- Significant spikes at lags 7, 14, and 21 (visible well above the confidence band).
- This pattern suggests seasonality (possibly weekly, every 7 days).
- Indicates potential values for q = 1, 2, or 3.

# PACF (Partial Autocorrelation Function):

- Significant spikes at lags 7 and 21.
- Gradual decline after lag 1.
- Indicates potential values for p = 1 or 2.

# Modeling

```
In [46]: def performance(actual, predicted):
    print('MAE :', round(mae(actual, predicted), 3))
    print('RMSE :', round(mse(actual, predicted)**0.5, 3))
    print('MAPE:', round(mape(actual, predicted), 3))
```

### **Arima Model**

```
In [47]: X_train, X_test = Stationary_data.iloc[:-35], Stationary_data.iloc[-35:]

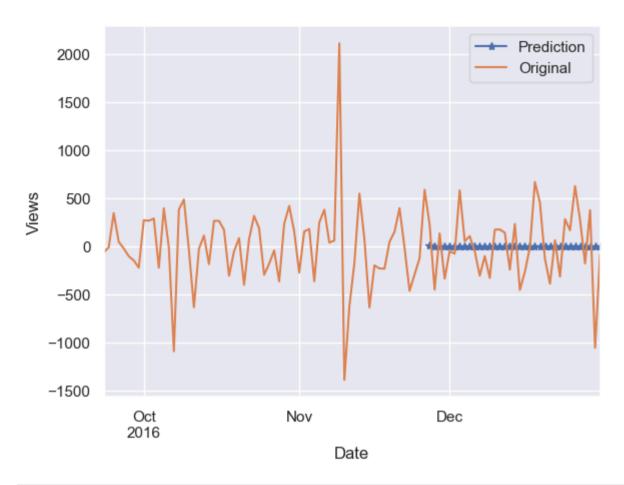
# Fit the ARIMA model
model = sm.tsa.ARIMA(endog=X_train, order=(1,1,1), seasonal_order=(0,0,0,0))
result = model.fit()

# Make predictions
predTest = result.forecast(len(X_test))

# Plot the predictions and the original data
predTest.plot(style="-*", label="Prediction")
Stationary_data.tail(100).plot(label="Original")
plt.ylabel("Views")
plt.xlabel("Date")
plt.legend()

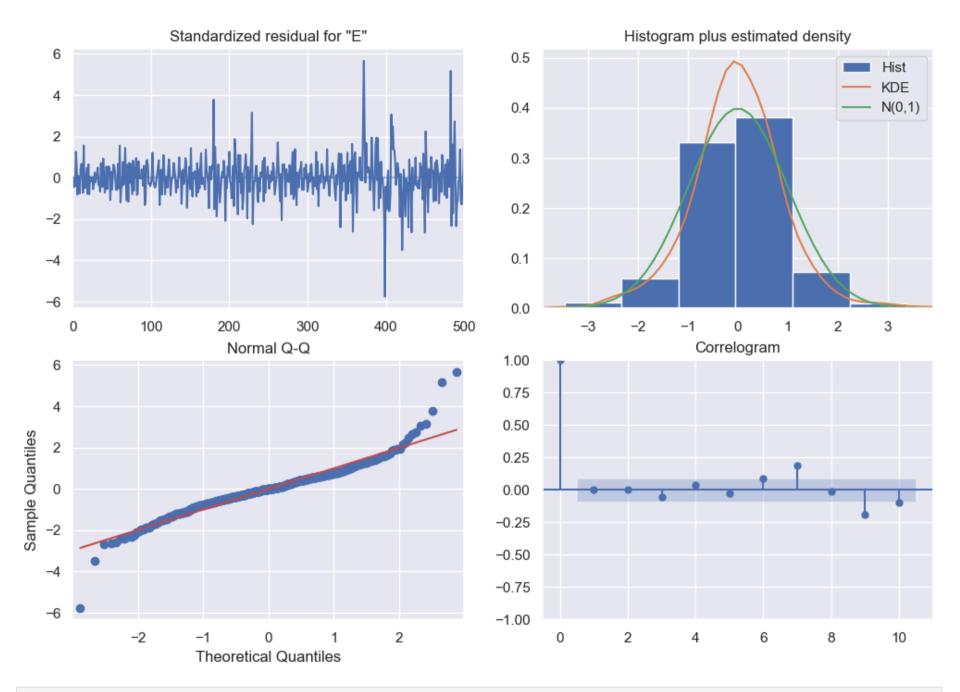
# Evaluate the performance of the model
performance(X_test, predTest)
print()
```

MAE: 272.444 RMSE: 347.019 MAPE: 1.006



```
# Fit the model
        fit = model.fit(disp=False)
       forecast = fit.forecast(steps=len(X test), exog=X test exog)
In [52]:
       print(fit.summary())
In [53]:
        # Plot diagnostics
        fit.plot diagnostics(figsize=(12, 8))
        plt.show()
                                      SARIMAX Results
        _____
       Dep. Variable:
                                          English
                                                  No. Observations:
                                                                               514
                       SARIMAX(1, 1, 1)x(1, 1, 1, 12)
       Model:
                                                  Log Likelihood
                                                                          -3714,472
       Date:
                                   Mon, 21 Apr 2025
                                                                           7440.944
                                                  AIC
       Time:
                                         12:11:44
                                                  BIC
                                                                           7466,243
        Sample:
                                               0
                                                  HOIC
                                                                           7450.871
                                            - 514
       Covariance Type:
                                             opg
        ______
                                                P>|z|
                                                         [0.025
                                                                   0.975]
                      coef
                            std err
                                          Z
        Exog
                 -1435.7917
                             57.935
                                     -24.783
                                                0.000
                                                       -1549.342
                                                                -1322.241
        ar.L1
                    0.2925
                              0.055
                                       5.320
                                                0.000
                                                          0.185
                                                                   0.400
                   -0.8304
                              0.033
                                     -24.873
                                                0.000
                                                         -0.896
                                                                   -0.765
       ma.L1
                                      -3.781
                                                         -0.256
                                                                   -0.081
        ar.S.L12
                   -0.1685
                              0.045
                                                0.000
       ma.S.L12
                   -0.9846
                              0.090
                                     -10.899
                                                0.000
                                                         -1.162
                                                                   -0.808
        sigma2
                 1.504e+05
                           1.29e+04
                                      11.651
                                                0.000
                                                       1.25e+05
                                                                 1.76e+05
        ______
       Ljung-Box (L1) (Q):
                                            Jarque-Bera (JB):
                                       0.00
                                                                       740.48
       Prob(Q):
                                            Prob(JB):
                                       0.96
                                                                        0.00
       Heteroskedasticity (H):
                                       4.62
                                            Skew:
                                                                        0.27
       Prob(H) (two-sided):
                                            Kurtosis:
                                                                        8.93
                                       0.00
       Warnings:
```

[1] Covariance matrix calculated using the outer product of gradients (complex-step).



```
RMSE: 575.243
         MAPE: 3.067
In [80]: X train, X test = English clipped.iloc[:-35], English clipped.iloc[-35:]
In [84]: X train exog, X test exog = exog.iloc[:-35], exog.iloc[-35:]
In [85]: X train.index = X train exog.index
In [55]: # Define parameter ranges
         p = range(0, 3) # AR order
         d = range(0, 2) # Differencing order
         q = range(0, 3) # MA order
         P = range(0, 3) # Seasonal AR order
         D = range(0, 2) # Seasonal differencing order
         Q = range(0, 3) # Seasonal MA order
         s = [12] # Seasonal period (e.g., monthly data)
         # Create all combinations of parameters
         param grid = list(itertools.product(p, d, q, P, D, Q, s))
In [56]: results = []
         # Iterate over parameter combinations
         for params in param grid:
             try:
                 # Define the SARIMAX model
                 model = SARIMAX(
                     X train,
                     order=(params[0], params[1], params[2]),
                     seasonal order=(params[3], params[4], params[5], params[6]),
                     exog=X train exog # Include exogenous variables if available
                 # Fit the model
                 fit = model.fit(disp=False)
                 # Save parameters and evaluation metric
                 results.append({
                     'params': params,
                     'aic': fit.aic,
                     'bic': fit.bic
```

MAE: 430.316


```
except Exception as e:
                 # Save failed parameter combination
                 results.append({
                      'params': params,
                      'aic': None,
                      'bic': None,
                      'error': str(e)
                 })
In [58]: results df = pd.DataFrame(results)
          # Drop rows with errors or NaN AIC
         results df = results df.dropna(subset=['aic'])
         # Sort by AIC
         best model = results df.sort values(by='aic').iloc[0]
          # Display the best parameters
         print("Best Parameters:", best model['params'])
         print("AIC:", best model['aic'])
         Best Parameters: (1, 0, 1, 1, 1, 2, 12)
         AIC: 14.0
         best params = best model['params']
In [86]:
         # Refit the model with best parameters
         final model = SARIMAX(
             X train,
              order=(best params[0], best params[1], best params[2]),
             seasonal order=(best params[3], best params[4], best params[5], best params[6]),
              exog=X train exog
         ).fit(disp=False)
         # Forecast using the final model
         forecast = final model.forecast(steps=len(X test), exog=X test exog)
In [61]: print(fit.summary())
         # Plot diagnostics
         fit.plot diagnostics(figsize=(12, 8))
         plt.show()
```

#### SARIMAX Results

Dep. Varia Model: Date:			2)x(2, 1, 2 Mon, 21 Apr	, 12) Log 2025 AIC	Observation Likelihood	======= S:	514 -3739.551 7499.101
Time:			12:	51:10 BIC			7541.267
Sample:				0 HQI	C		7515.646
				- 514			
Covariance	e Type:			opg			
=======	coef				[0.025	_	
Exog	-1453.9972		-15.819		-1634.145		
ar.L1	-0.7339	0.307	-2.394	0.017	-1.335	-0.133	
ar.L2	0.2659	0.116	2.283	0.022	0.038	0.494	
ma.L1	0.1957	0.552	0.354	0.723	-0.887	1.278	
ma.L2	-0.8035	0.455	-1.767	0.077	-1.694	0.088	
ar.S.L12	0.8142	0.215	3.789	0.000	0.393	1.235	
ar.S.L24	0.1643	0.085	1.931	0.053	-0.002	0.331	
ma.S.L12	-1.9790	0.475	-4.169	0.000	-2.909	-1.049	
ma.S.L24	0.9804	0.469	2.089	0.037	0.060	1.900	
sigma2	2.399e+05	0.006	4.02e+07	0.000	2.4e+05	2.4e+05	
Ljung-Box	(L1) (Q):	:=======	.====== 0.03	Jarque-Bera	======= a (JB):	 74	:==== !3.56
Prob(0):			0.86	•	` '		0.00
Heterosked	lasticity (H):		4.64	• •			0.28
	:wo-sided):		0.00	Kurtosis:			8.94
=======	========		=======	=======	=======	========	====

#### Warnings:

- [1] Covariance matrix calculated using the outer product of gradients (complex-step).
- [2] Covariance matrix is singular or near-singular, with condition number 8.48e+24. Standard errors may be unstable.

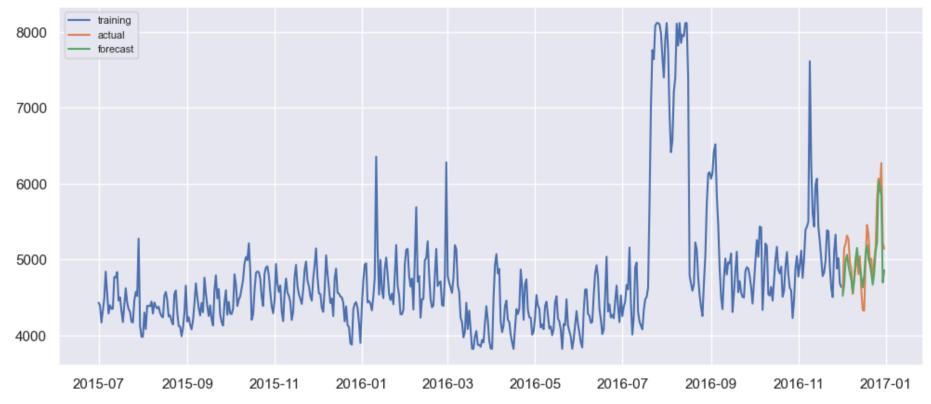


MAE : 396.736 RMSE : 473.142 MAPE: 0.076

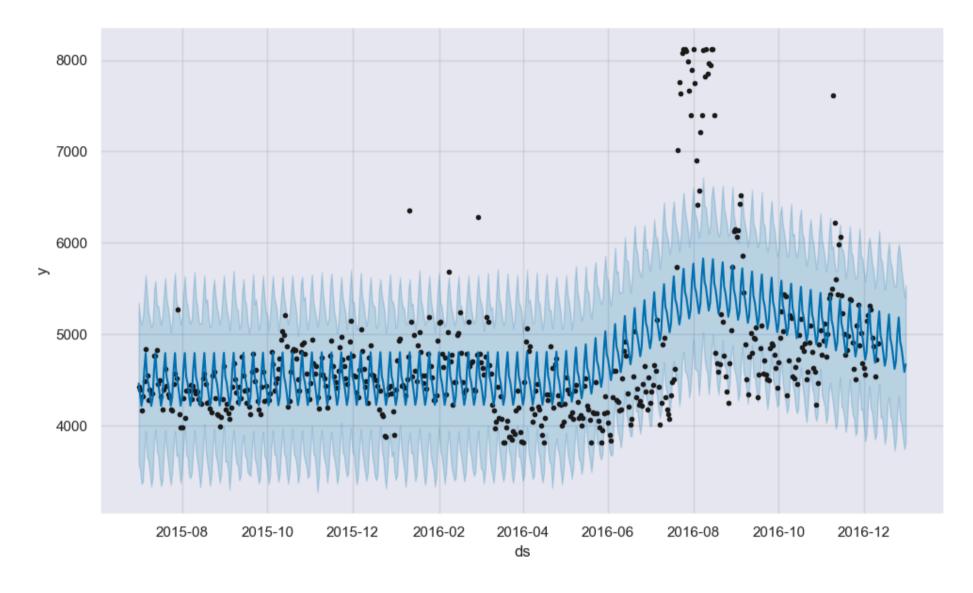
# Predicting the future Values using SARIMAX

```
In [63]: ex=exog['Exog'].to_numpy()
In [88]: train=English_clipped[:520]
         test=English clipped[520:]
         model=sm.tsa.statespace.SARIMAX(train,order=(4, 1, 3),seasonal_order=(1,1,1,7),exog=ex[:520])
         results=model.fit()
         fc=results.forecast(30,dynamic=True,exog=pd.DataFrame(ex[520:]))
         # Make as pandas series
         fc series = pd.Series(fc)
          # Plot
         train.index=train.index.astype('datetime64[ns]')
         test.index=test.index.astype('datetime64[ns]')
         plt.figure(figsize=(12,5), dpi=100)
         plt.plot(train, label='training')
         plt.plot(test, label='actual')
         plt.plot(fc series, label='forecast')
         plt.title('Forecast vs Actuals')
         plt.legend(loc='upper left', fontsize=8)
         <matplotlib.legend.Legend at 0x206be5b5450>
Out[88]:
```

#### Forecast vs Actuals



# **FB Prophet**



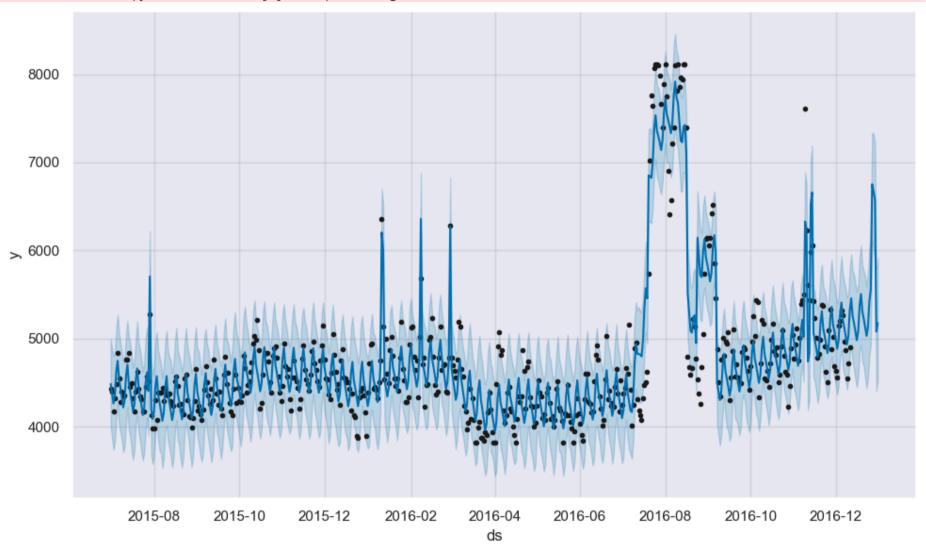
# Running FB Prophet using Exogeneous Variable

```
In [101... df['exog'] = exog

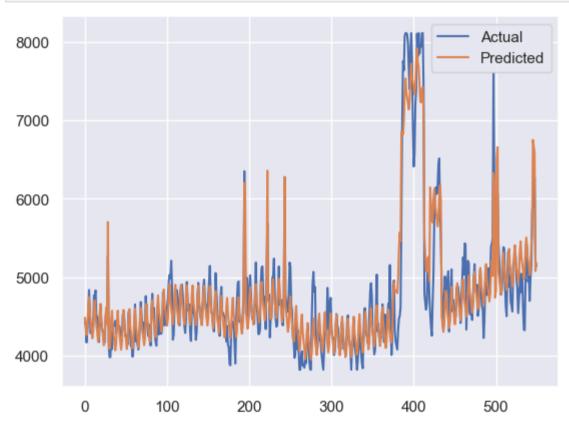
In [102... model2=Prophet(interval_width=0.9, weekly_seasonality=True, changepoint_prior_scale=1)
    model2.add_regressor('exog')
    model2.fit(df[:-20])
```

```
forecast2 = model2.predict(df)
fig = model2.plot(forecast2)
```

```
13:13:05 - cmdstanpy - INFO - Chain [1] start processing 13:13:06 - cmdstanpy - INFO - Chain [1] done processing
```



```
plt.plot(y_pred, label='Predicted')
plt.legend()
plt.show()
```



```
In [105... mape = np.mean(np.abs(forecast2['yhat'][-20:] - df['y'][-20:].values)/np.abs(df['y'][-20:].values))
print("mape:",mape)
```

mape: 0.06620436903062904

- Fb prophet 6.6% Mape
- Sarimax 7.6% Mape

# Recommendations based on MAPE & mean\_visits:

- English language is a clear winner. Maximum advertisement should be done on English pages. Their MAPE is low & mean visits are high.
- Chinese language has lowest number of visits. Advertisements on these pages should be avoided unless business has specific marketing strategy for Chinese populations.
- Russian language pages have decent number of visits and low MAPE. If used properly, these pages can result in maximum conversion.
- Spanish language has second highest number of visits but their MAPE is highest. There is a possibility advertisements on these pages won't reach the final people.
- French, German & Japenese have medium level of visits & medium MAPE levels.
- Depending on target customers advertisements should be run on these pages.

# Questions

### Defining the problem statements and where can this and modifications of this be used?

• We are provided with the data of 145k wikipedia pages and daily view count for each of them. Our clients belong to different regions and need data on how their ads will perform on pages in different languages. By creating a proper forecasting model to predict the fluctuations of visits on pages, we can help the business team to optimise the marketing spend. If we can predict the days with higher visits properly, the business will run the ads for those specific days and still be able to reach wider audience with most optimized spend.

### Write 3 inferences you made from the data visualizations

- All-agents access is significantly higher than spider access across all languages, with English having the highest overall count.
- Japanese and German also show high engagement from all-agents, while Spanish and Russian have the lowest counts in both categories.
- English has more pages with AccessType of all-access, different from the rest
- English pages are the most visited pages follwed by Spanish
- English pages have an upward trend in terms of visits
- There is an unusual peak from mid of July to end of August 2016

## What does the decomposition of series do?

• The decomposition of a time series is a technique used to separate the series into its constituent components.

### What level of differencing gave you a stationary series?

• Level 1 (diff(1))

### Difference between arima, sarima & sarimax

- ARIMA:
  - Models non-seasonal time series.
  - No seasonality or external factors.
- SARIMA:
  - Extends ARIMA to handle seasonality.
  - No external factors.
- SARIMAX:
  - Extends SARIMA by including external predictors (exogenous variables).
  - Handles seasonality and external factors.

### Compare the number of views in different languages

• Mean number of views (Popularity sequence) of various languages have the following: English > Spanish > Russian > German > Japenese > French > Chinese

# What other methods other than grid search would be suitable to get the model for all languages?

- Bayesian Optimization
- Randomized Search